

# **CSE 4020/5260**

## **Database Systems**

Instructor: Fitzroy Nembhard, Ph.D.

**Week 8 - 9**

# **DML**



# Distribution

---

- All slides included in this class are for the exclusive use of students and instructors associated with Database Systems (CSE 4020/5260) at the Florida Institute of Technology
- Redistribution of the slides is not permitted without the written consent of the author.

# Structured Query Language (SQL)

---

- Basic SQL Query Structure
- Set Operations
- Aggregate Functions
- Nested Subqueries
- Derived Relations
- Views
- Modification of a Database
- Specialized Join Operation

# Banking Example

---

- SQL is a “standardized” language, but most vendors have their own version.
- Queries are typically submitted on the command-line, using a client query tool, via program code, or through an API.
- Now is the time to start issuing queries, just to get the hang of it!
- White space will be used liberally throughout.

# Banking Example

- Recall the banking database:

*branch* (*branch-name*, *branch-city*, *assets*)

*customer* (*customer-name*, *customer-street*, *customer-city*)

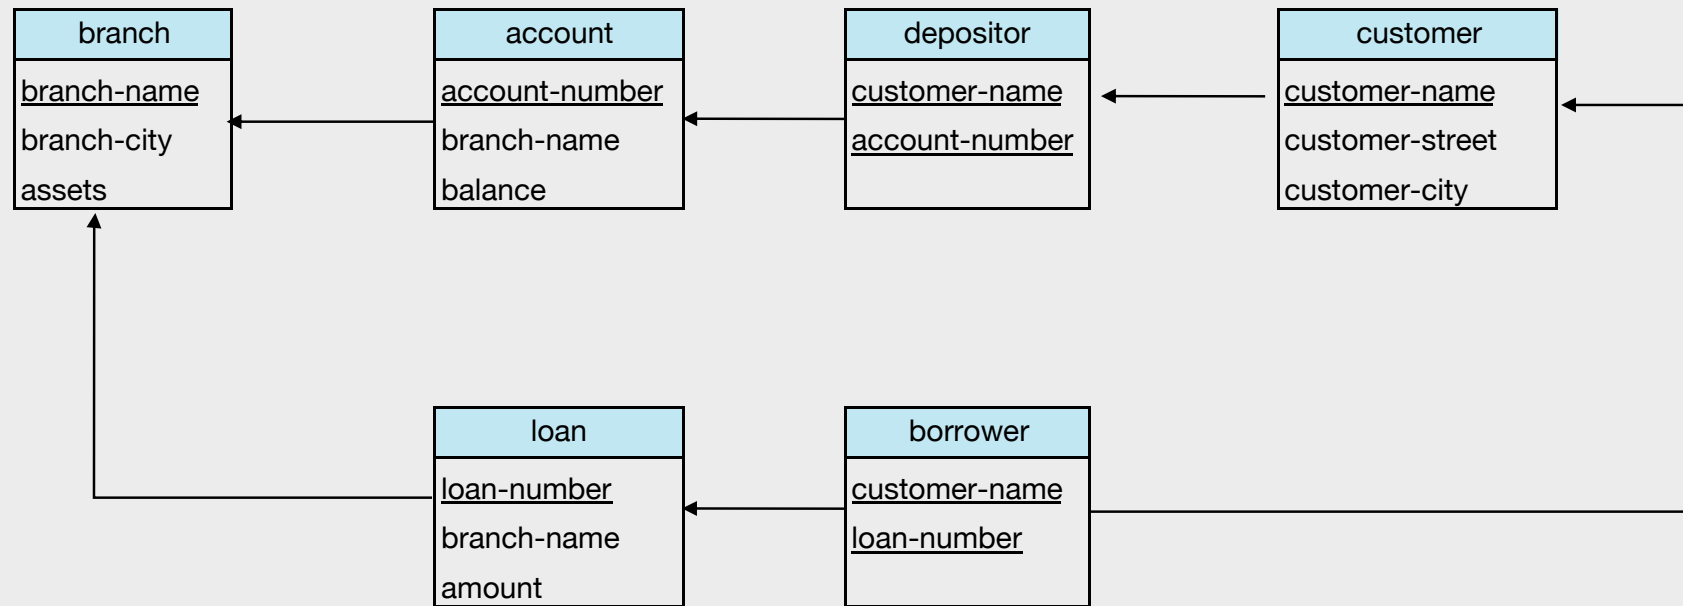
*account* (*account-number*, *branch-name*, *balance*)

*loan* (*loan-number*, *branch-name*, *amount*)

*depositor* (*customer-name*, *account-number*)

*borrower* (*customer-name*, *loan-number*)

# Schema Used in Examples



# Basic Structure

- Typical SQL statement/query structure:

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_m$   
where  $P$ 
```

- Equivalent (sort of) to:

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$



# The select Clause

- **select** clause - lists desired attributes (corresponds to projection).

*“Find the names of those branches that have outstanding loans.”*

```
select branch-name  
from loan
```

$$\Pi_{\text{branch-name}}(\text{loan})$$

```
select branch-name, loan-number  
from loan
```

$$\Pi_{\text{branch-name, loan-number}}(\text{loan})$$

## Schema

branch (branch-name, branch-city, assets)

customer (customer-name, customer-street, customer-city)

account (account-number, branch-name, balance)

loan (loan-number, branch-name, amount)

depositor (customer-name, account-number)

borrower (customer-name, loan-number)



# The select Clause (Cont.)

- An asterisk denotes all attributes:

```
select *  
from loan
```

- **select** can contain expressions (corresponds to generalized projection).

```
select loan-number, branch-name, amount * 100  
from loan
```

- Note that the above does not modify the table.

Generalized projection:

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$



## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# The select Clause (Cont.)

- The basic SQL select statement does NOT eliminate duplicates.
- Keyword **distinct** is used to eliminate duplicates.

*“Find the names of those branches that have outstanding loans (no duplication).”*

```
select distinct branch-name  
from loan
```

- Keyword **all** can be used (redundantly) when duplicates desired.

```
select all branch-name  
from loan
```

## Schema

branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)

# The where Clause

- **where** clause - specifies conditions on the result (corresponds to selection).

*“Find the loan numbers for all loans over \$1200 made at the Perryridge branch.”*

```
select loan-number
      from loan
      where branch-name = 'Perryridge' and amount > 1200
```

SELECT column1, column2, ...  
FROM table\_name  
WHERE **NOT** condition;



- Logical connectives **and**, **or**, and **not** can be used.
- Comparisons can be applied to results of arithmetic expressions.

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# The from Clause

- **from** clause - lists required relations (corresponds to Cartesian product).

*“Find the Cartesian product borrower x loan.”*

```
select * from borrower, loan
```

*“Find the name, loan number and loan amount for all customers having a loan at the Perryridge branch.”*

```
select borrower.customer-name, borrower.loan-number, loan.amount  
from borrower, loan  
where borrower.loan-number = loan.loan-number and  
loan.branch-name = 'Perryridge'
```

- Note the use of expanded name notation in the above.

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# The from Clause

- Sometimes mixed-use notation is used:

```
select customer-name, borrower.loan-number, amount  
from borrower, loan  
where borrower.loan-number = loan.loan-number and  
       branch-name = 'Perryridge'
```

## Schema

```
branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)
```

# The Rename Operation

- Attribute renaming (**as**):
- In the **select** clause (for column renaming):

*“Find the name, loan number and loan amount of all customers; rename the loan-number column loan-id.”*

```
select customer-name, borrower.loan-number as loan-id, amount
from borrower, loan
where borrower.loan-number = loan.loan-number
```

## Schema

branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)

# Tuple Variables

- In the **from** clause (for abbreviating):

*“Find the customer names, their loan numbers and loan amounts for all customers having a loan at the Perryridge branch.”*

```
select T.customer-name, T.loan-number, S.amount  
       from borrower as T, loan as S  
       where T.loan-number = S.loan-number  
       and S.branch-name = 'Perryridge'
```

## Schema

branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)



# Tuple Variables

- It can also be used to resolve ambiguous relation names:

*“Find the names of all branches that have greater assets than some branch located in Brooklyn.”*

```
select distinct T.branch-name  
  from branch as T, branch as S  
  where T.assets > S.assets and S.branch-city = 'Brooklyn'
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# String Operations

- So how about strings?
- SQL supports a variety of string processing functions...surprise!!!
- Example:

*“Find the names of all customers whose street includes the substring ‘Main’.”*

```
select customer-name  
      from customer  
      where customer-street like '%Main%'
```

# String Operations

- Other SQL string operations:
  - concatenation (using “||”)
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.
  
- Most COTS DBMS query processors augment SQL string processing with even more operations; the list is typically very long.

# Ordering the Display of Tuples

- Sorting:

*“List in alphabetic order the names of all customers having a loan at the Perryridge branch.”*

```
select distinct customer-name
from borrower, loan
where borrower.loan-number = loan.loan-number and
       branch-name = 'Perryridge'
order by customer-name
```

- **desc** or **asc** (the default) can be specified:

- **order by** *customer-name desc*

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Ordering the Display of Tuples

- Sorting on multiple attributes (with both **asc** and **desc**):
- Example: add loan amount to the previous query:

```
select distinct customer-name, amount  
from borrower, loan  
where borrower.loan-number = loan.loan-number and  
       branch-name = 'Perryridge'  
order by customer-name asc, amount desc
```

## Schema

```
branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)
```

# Set Operations

- **union, intersect, and except** ( $\cup$ ,  $\cap$ ,  $-$ , respectively):

- $r \text{ union } s$
- $r \text{ intersect } s$
- $r \text{ except } s$

where  $r$  and  $s$  are either relations or sub-queries.

- The above operations all automatically eliminate duplicates.

- Note: MySQL does not support the **except** clause (unlike Oracle, which uses the minus keyword in place of except, MySQL does not support except at all). You can use the **not in** clause instead.

# Set Operations

*“Find all customers who have a loan, an account, or both.”*

```
(select customer-name from depositor)
union
(select customer-name from borrower)
```

*“Find all customers who have both a loan and an account.”*

```
(select customer-name from depositor)
intersect
(select customer-name from borrower)
```

*“Find all customers who have an account but no loan.”*

```
(select customer-name from depositor)
except
(select customer-name from borrower)
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)



# Set Operations

- **union all**, **intersect all** and **except all** retain duplicates:

If a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:

- $m + n$  times in  $r$  **union all**  $s$
- $\min(m, n)$  times in  $r$  **intersect all**  $s$
- $\max(0, m - n)$  times in  $r$  **except all**  $s$

# Aggregate Functions

- Grouping and aggregate functions.

- Basic aggregate functions:

<b>avg</b>	- average value
<b>min</b>	- minimum value
<b>max</b>	- maximum value
<b>sum</b>	- sum of values
<b>count</b>	- number of values

- Aggregate functions operate on groups.

# Aggregate Functions, Cont.

*“Find the average account balance.”*

```
select avg (balance)  
from account
```

*“Find the average account balance at the Perryridge branch.”*

```
select avg (balance)  
from account  
where branch-name = 'Perryridge'
```

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Aggregate Functions, Cont.

*“Find the number of tuples in the depositor relation.”*

```
select count (*)  
from depositor
```

Or any single or combination of columns:

```
select count (customer-name)  
from depositor
```

```
select count (account-number)  
from depositor
```

```
select count (customer-name, account-number)  
from depositor
```

## Schema

```
branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)
```

# Aggregate Functions, Cont.

*“Find the number of depositors in the bank.”*

```
select count (distinct customer-name)  
  from depositor
```

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Aggregate Functions – Group By

- Aggregate functions applied to groups:

*“Find the number of accounts for each branch.”*

```
select branch-name, count (account-number)
from account
group by branch-name
```

*“Find the number of depositors for each branch.”*

```
select branch-name, count (distinct customer-name)
from depositor, account
where depositor.account-number = account.account-number
group by branch-name
```

- Why does the second have *distinct* but not the first?

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# Aggregate Functions – Group By

- Grouping can be on multiple attributes:

*“For each depositor, determine how many accounts that depositor has at each branch.”*

```
select customer-name, branch-name, count (depositor.account-number)
from depositor, account
where depositor.account-number = account.account-number
group by customer-name, branch-name
```

- Notes:

- Should *distinct* have been included?
- Attributes in the **select** clause outside of the aggregate functions must appear in **group by** list (e.g., delete *branch-name* from the group-by clause).
- Group-by *might* require a sort.

## Schema

*branch* (branch-name, branch-city, assets)

*customer* (customer-name, customer-street, customer-city)

*account* (account-number, branch-name, balance)

*loan* (loan-number, branch-name, amount)

*depositor* (customer-name, account-number)

*borrower* (customer-name, loan-number)



# Aggregate Functions – Group By

- Grouping on multiple attributes, and multiple aggregate functions.

*“For each depositor, determine how many accounts that depositor has at each branch, plus the average, min and max balance for any account at that branch.”*

```
select customer-name,  
        branch-name,  
        count (depositor.account-number)  
        avg (account.balance)  
        min (account.balance)  
        max (account.balance)  
from depositor, account  
where depositor.account-number = account.account-number  
group by customer-name, branch-name
```

## Schema

```
branch (branch-name, branch-city, assets)  
customer (customer-name, customer-street, customer-city)  
account (account-number, branch-name, balance)  
loan (loan-number, branch-name, amount)  
depositor (customer-name, account-number)  
borrower (customer-name, loan-number)
```

# Aggregate Functions – Having Clause

- Groups can be selected or eliminated using the *having* clause.

*“Find those branches in Orlando with an average balance over 1200.”*

```
select branch-name
  from account, branch
 where account.branch-name = branch.branch-name
       and branch-city = 'Orlando'
 group by branch-name
 having avg (balance) > 1200
```

- Predicates in the **having** clause are applied after the formation of groups, but those in the **where** clause are applied before forming groups.

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# Null Values

- It is possible for tuples to have a *null* value for some attributes.
- *null* signifies an unknown value or that a value does not exist.
- The rules for null values are consistent with relational algebra (repeated on the following pages), except for the following addition...
- The predicate **is null** can be used to check for null values.

*“Find all loan numbers in the loan relation with null values for amount.”*

```
select loan-number  
from loan  
where amount is null
```

# Null Values and Three Valued Logic

- Rule #1 - Any comparison with *null* (initially) returns *unknown*:

➤  $5 < null$  or  $null < null$  or  $null = null$

```
select loan-number
from loan
where amount > 50
```

```
select borrower-name, branch-name
from borrower, loan
where borrower.loan-number = loan.loan-number
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

- Rule #2 - The result of any arithmetic expression involving *null* is *null*

➤  $5 + null$  evaluates to *null*

```
select loan-number
from loan
where amount*100 > 50000
```

# Null Values and Three Valued Logic

## ■ Rule #3 - A “three-valued logic” is applied to complex expressions:

- OR: (*unknown or true*) = *true*, (*unknown or false*) = *unknown* (*unknown or unknown*) = *unknown*
- AND: (*true and unknown*) = *unknown*, (*false and unknown*) = *false*, (*unknown and unknown*) = *unknown*
- NOT: (**not** *unknown*) = *unknown*
- “*P is unknown*” evaluates to *true* if predicate *P* evaluates to *unknown*

```
select loan-number
from loan
where amount*100 > 5000 and branch-name = “Perryridge”
```

## ■ Rule #4 - Final result of a **where** clause predicate is treated as *false* if it evaluates to *unknown*.

```
select loan-number
from loan
where amount*100 > 5000 and branch-name = “Perryridge”
```

# Null Values, Cont.

- Rule #5 - Aggregate functions, except **count**, simply ignore nulls.
- Total all loan amounts:

```
select sum (amount)  
from loan
```

- Above statement ignores null amounts
- Result is null if there is no non-null amount

# Null Values and Expression Evaluation, Cont.

- This all seems like a pain...couldn't it be simplified?
- Why doesn't a comparison with *null* simply result in *false*?

If *false* was used instead of *unknown*, then:

$\text{not } (A < 5)$

would not be equivalent to:

$A \geq 5$

Why would this be a problem?



# Nested Subqueries

---

- SQL provides a mechanism for nesting queries.
- A sub-query is a **select** statement that is nested in another SQL query.
- Nesting is usually in a **where** clause but may be in a **from** clause.

# Nested Subqueries

- Sub-query in a **where** clause typically performs a set test.

in	<comp> some	exists	unique
not in	<comp> all	not exists	not unique

where <comp> can be <, ≤, >, =, ≠

# Example Query

*“Find all customers who have both an account and a loan.”*

```
select distinct customer-name  
from borrower  
where customer-name in (select customer-name  
                           from depositor)
```

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Example Query

*“Find all customers who have a loan but do not have an account.”*

```
select distinct customer-name  
from borrower  
where customer-name not in (select customer-name  
                                from depositor)
```

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Example Query

*“Find the names of all customers who have both an account and a loan at the Perryridge branch.”*

```
select distinct customer-name
from borrower, loan
where borrower.loan-number = loan.loan-number and
       branch-name = “Perryridge” and
       (branch-name, customer-name) in
       (select branch-name, customer-name
        from depositor, account
        where depositor.account-number =
              account.account-number)
```

=> Note that the above query can be “simplified.”

# Example Query

*“Find the names of all customers who have both an account and a loan at the Perryridge branch.”*

```
select distinct customer-name
from borrower, loan
where borrower.loan-number = loan.loan-number and
       branch-name = “Perryridge” and
       customer-name in
           (select customer-name
            from depositor, account
            where depositor.account-number =
                  account.account-number and
                  branch-name = “Perryridge”)
```

```
select distinct customer-name
from borrower, loan
where borrower.loan-number = loan.loan-number and
       branch-name = “Perryridge” and
       (branch-name, customer-name) in
           (select branch-name, customer-name
            from depositor, account
            where depositor.account-number =
                  account.account-number)
```

# Set Comparison – the “Some” Clause

*“Find all branches that have greater assets than some branch located in Brooklyn.”*

```
select distinct T.branch-name  
  from branch as T, branch as S  
  where T.assets > S.assets and  
        S.branch-city = 'Brooklyn'
```

Same query using > **some** clause:

```
select branch-name  
  from branch  
  where assets > some  
              (select assets  
                from branch  
                where branch-city = 'Brooklyn')
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# Set Comparison – the “All” Clause

*“Find the names of all branches that have greater assets than all branches located in Brooklyn.”*

```
select branch-name
from branch
where assets > all
      (select assets
       from branch
       where branch-city = 'Brooklyn')
```

Note that the some and all clauses correspond to existential and universal quantification, respectively.

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)



# Definition of the “Some” Clause

■  $F \text{ <comp> some } r \Leftrightarrow \exists t \in r \text{ s.t. } (F \text{ &ltcomp> } t)$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$

# Definition of the “All” Clause

■  $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{all}) \equiv \text{not in}$

However,  $(= \text{all}) \not\equiv \text{in}$

# Test for Empty Relations

- The **exists** operator can be used to test if a relation is empty.
- Operator **exists** returns **true** if its argument is nonempty.
  - **exists**  $r \Leftrightarrow r \neq \emptyset$
  - **not exists**  $r \Leftrightarrow r = \emptyset$
- On a personal note, why not call it **empty**?

# Example Query

“Find all customers who have an account at all branches located in Brooklyn.”

```
select distinct S.customer-name
  from customer as S
 where not exists (
    (select branch-name
     from branch
     where branch-city = 'Brooklyn')
    except
    (select R.branch-name
     from depositor as T, account as R
     where T.account-number = R.account-number and
           S.customer-name = T.customer-name))
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

- Because of the use of the tuple variable S in the nested query, the above is sometimes referred to as a *correlated* query.
- The above demonstrates that nesting can be almost arbitrarily composed and deep.
- According to the book, the above cannot be written using = **all** or its variants...hmmm...
- Recall that except does not work in MySQL

# Test for Absence of Duplicate Tuples

- The **unique** operator tests whether a sub-query contains duplicate tuples.

*“Find all customers who have at most one account at the Perryridge branch.”*

```
select T.customer-name
  from customer as T
  where unique (
    select D.customer-name
    from account as A, depositor as D
    where T.customer-name = D.customer-name and
          A.account-number = D.account-number and
          A.branch-name = 'Perryridge')
```

- What if the inner query selected the account number?
  - `count(...) <= 1`

# Example Query

*“Find all customers who have at least two accounts at the Perryridge branch.”*

```
select distinct T.customer-name
from customer T
where not unique (
    select R.customer-name
from account, depositor as R
where T.customer-name = R.customer-name and
       R.account-number = account.account-number and
       account.branch-name = 'Perryridge')
```

## Schema

*branch* (*branch-name*, *branch-city*, *assets*)  
*customer* (*customer-name*, *customer-street*, *customer-city*)  
*account* (*account-number*, *branch-name*, *balance*)  
*loan* (*loan-number*, *branch-name*, *amount*)  
*depositor* (*customer-name*, *account-number*)  
*borrower* (*customer-name*, *loan-number*)

# Nesting in the From-Clause

*“Find the average account balance of those branches where the average account balance is greater than \$1200.”*

```
select branch-name, avg-balance
  from (select branch-name, avg (balance)
        from account
        group by branch-name)
      as result (branch-name, avg-balance)
 where avg-balance > 1200
```

Note that previously we saw an equivalent query that used a *having* clause.

# Views

## ■ Purpose of a view:

- Hide certain data from the view of certain users
- Provide pre-canned, named queries
- Simplify complex queries

## ■ Syntax of a view:

**create view *v* as** <query expression>

where:

- *v* - view name
- <query expression> - view definition (SQL)



# Example Views

- A view consisting of branches and their customers:

```
create view all-customer as
  (select branch-name, customer-name
   from depositor as D, account as A
   where D.account-number = A.account-number)
  union
  (select branch-name, customer-name
   from borrower as B, loan as L
   where B.loan-number = L.loan-number)
```

*“Find all customers of the Perryridge branch.”*

```
select customer-name
  from all-customer
 where branch-name = 'Perryridge'
```

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# Modification of the Database – Insertion

## ■ Basic insert:

```
insert into account values ('A-9732', 'Perryridge', 1200)
```

## ■ Ordering values:

```
insert into account (branch-name, balance, account-number)
values ('Perryridge', 1200, 'A-9732')
```

## ■ Inserting a null value:

```
insert into account values ('A-777', 'Perryridge', null)
```

### Schema

```
branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)
```

# Modification of the Database – Insertion

*“Provide as a gift for all loan customers of the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new account.”*

```
insert into account
  select loan-number, branch-name, 200
  from loan
  where branch-name = 'Perryridge'
```

```
insert into depositor
  select customer-name, loan-number
  from loan, borrower
  where branch-name = 'Perryridge'
        and loan.account-number = borrower.account-number
```

- The above would typically be a transaction.

## Schema

*branch* (branch-name, branch-city, assets)  
*customer* (customer-name, customer-street, customer-city)  
*account* (account-number, branch-name, balance)  
*loan* (loan-number, branch-name, amount)  
*depositor* (customer-name, account-number)  
*borrower* (customer-name, loan-number)

# Modification of the Database – Insertion

- Most DBMSs provide a command-line, bulk-load command:

```
LOAD DATA LOCAL INFILE '<file-path>' INTO TABLE part
    FIELDS TERMINATED BY '<file-separator>'
    LINES TERMINATED BY '<line-separator>';
```

- Example:

```
LOAD DATA LOCAL INFILE '../drfitz/department.csv' INTO TABLE department
    FIELDS TERMINATED BY ',' -- Specify the delimiter
    ENCLOSED BY '"' -- When strings are enclosed in quotes
    LINES TERMINATED BY '\n'
    IGNORE 1 LINES -- or ROWS
```

# Modification of the Database – Deletion

*“Delete all tuples in the depositor table.”*

**delete from** *depositor*

*“Delete all depositor records for Smith.”*

**delete from** *depositor*  
**where** *customer-name* = 'Smith'

# Modification of the Database – Deletion

*“Delete all accounts at every branch located in Needham city.”*

```
delete from depositor
where account-number in
    (select account-number
     from branch as B, account
     where branch-city = 'Needham' and B.branch-name = A.branch-name)
```

```
delete from account
where branch-name in (select branch-name
                        from branch
                        where branch-city = 'Needham')
```

# Example Query

*“Delete the record of all accounts with balances below the average at the bank.”*

```
delete from account  
where balance < (select avg (balance)  
                  from account)
```

# Modification of the Database – Updates

*“Set the balance of all accounts at the Perryridge branch to 0.”*

```
update account  
  set balance = 0  
  where branch-name = “Perryridge”
```

*“Set the balance of account A-325 to 0, and also change the branch name to “Mianus.”*

```
update account  
  set balance = 0, branch-name = “Mianus”  
  where account-number = “A-325”
```



# Modification of the Database – Updates

*“Increase all accounts with balances over \$10,000 by 6%, all other accounts by 5%.”*

## Option #1:

```
update account  
set balance = balance * 1.06  
where balance > 10000
```

```
update account  
set balance = balance * 1.05  
where balance <= 10000
```

# Modification of the Database – Updates

*“Increase all accounts with balances over \$10,000 by 6%, all other accounts by 5%.”*

Option #2:

```
update account
set balance = case
    when balance <= 10000 then balance * 1.05
    else balance * 1.06
end
```

# Transactions

- Some of the previous multi-query operations should be made transactions.
- A transaction is a sequence of SQL statements executed as a single unit.
- Example - Transferring money from one account to another:
  - deducting the money from one account
  - crediting the money to another account
- If one step succeeds and the other fails, the database is left in an inconsistent state.
- Therefore, either both steps should succeed, or both should fail (note: failing is better than corrupting).

# Transaction - Syntax

- Transactions are started either implicitly or explicitly.
- Transactions are terminated by:
  - *commit* - makes all updates of the transaction permanent
  - *rollback* - undoes all updates performed by the transaction
- Commits and rollbacks can also be either implicit or explicit.
- Implicit transactions with implicit commits (no special syntax):
  - DDL statements
  - Individual SQL statements that execute successfully
- Implicit rollbacks:
  - System failure

# Transactions, Cont.

- Automatic commit can be turned off, allowing multi-statement transactions.

- Transactions are identified by some variant of:

```
begin transaction           // shuts off auto-commit
...
end transaction             // commits the transaction
```

- Within the transaction, partial work can be:
  - made permanent by using the **commit work** statement.
  - undone by using the **rollback work** statement.
- Transactions are, or rather, should be the *rule* for programmers, rather than the exception.

# Which MYSQL Package should I Use?

- With which package should I code?

<b>MySQLdb</b>	A thin python wrapper around C module which implements API for MySQL database
<b>mysqlclient</b>	By far the fastest MySQL connector for CPython. Requires the mysql-connector-c C library to work.
<b>PyMySQL</b>	Pure Python MySQL client. Use if you can't use libmysqlclient
<b>mysql-connector-python</b>	MySQL connector developed by the MySQL group at Oracle

Recommended Python Adapters	
Database	Adapter
PostgreSQL	<a href="#">Psycopg</a>
SQLite	<a href="#">sqlite3</a>
Oracle	<a href="#">cx_oracle</a>
MySql	<a href="#">mysql-connector</a>

# Commit/Rollback Python Example

```
import mysql.connector
try:
    # Connecting to the Database
    conn = mysql.connector.connect(
        host='localhost',
        database=studentdb',
        user='root',
    )
    conn.autocommit = False # turn off autocommit to enforce transaction
    cursor = conn.cursor()
    statement = "UPDATE student SET dept_name ='Comp Sci.' WHERE id ='901000000'"
    cursor.execute(statement)

    # commit changes to the database
    conn.commit()

    # update successful message
    print("Student Department updated")

except mysql.connector.Error as error:
    conn.rollback()
conn.close()
```

# Commit/Rollback Python Example

```
import mysql.connector
try:
    # Connecting to the Database
    conn = mysql.connector.connect(
        host='localhost',
        database=studentdb',
        user='root',
        autocommit=True, #The autocommit flag may go here as well

    )
    conn.start_transaction()
    cursor = conn.cursor()
    # these two INSERT statements are executed as a single unit
    sql1 = """ insert into employees(name, salary) value('Jason', 75000) """
    sql2 = """ insert into employees(name, salary) value('Devon', 87000) """
    cursor.execute(sql1)
    cursor.execute(sql2)

    # commit changes to the database
    conn.commit()

    # display successful message
    print("Transaction committed")

except mysql.connector.Error as error:
    conn.rollback()
conn.close()
```



# Joined Relations

- Join operations take two relations and return another as a result.
- Specialized join operations are typically used as subquery expressions.
- Join condition – defines which tuples in the two relations match, and what attributes are present in the result of the join.
  - **natural**
  - **using** ( $A_1, A_2, \dots, A_n$ )      // equi-join
  - **on** <predicate>                      // theta-join
- Join type – defines how non-matching tuples (based on the join condition) in each relation are treated.
  - **inner join**
  - **left outer join**
  - **right outer join**
  - **full outer join**

# Joined Relations – Datasets for Examples

## ■ Relation *loan*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

## ■ Relation *borrower*

<i>customer-name</i>	<i>loan-number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

- Note that borrower information is missing for L-260 and loan information missing for L-155.

# Joined Relations – Examples

*loan inner join borrower*  
**on** *loan.loan-number = borrower.loan-number*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

loan

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

borrower

<i>customer-name</i>	<i>loan-number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

*loan left outer join borrower*  
**on** *loan.loan-number = borrower.loan-number*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>	<i>loan-number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	<i>null</i>	<i>null</i>

# Joined Relations – Examples

loan natural inner join borrower

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

loan natural right outer join borrower

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	null	null	Hayes

loan

loan-number	branch-name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

borrower

customer-name	loan-number
Jones	L-170
Smith	L-230
Hayes	L-155

## Joined Relations – Examples

*loan full outer join borrower using (loan-number)*

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>	<i>customer-name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes

*“Find all customers who have either an account or a loan (but not both) at the bank.”*

**select** *customer-name*  
**from** (*depositor natural full outer join borrower*)  
**where** *account-number is null or loan-number is null*

# Joined Relations – Examples

*“Find all customers who have an account at all branches located in Brooklyn.”*

```
select customer_name
from customer
where customer_name not in (
    select R.customer_name
    from depositor R
    inner join branch T on T.branch_city = 'brooklyn'
    left join account S
        on R.account_number = S.account_number
        and T.branch_name = S.branch_name
    where a.account_number is null )
```

## Schema

branch (branch-name, branch-city, assets)

customer (customer-name, customer-street, customer-city)

account (account-number, branch-name, balance)

loan (loan-number, branch-name, amount)

depositor (customer-name, account-number)

borrower (customer-name, loan-number)

# Exporting Data - Bonus

## Exporting a Table

```
SELECT * INTO OUTFILE 'data.txt'  
FIELDS TERMINATED BY ','  
FROM table2;
```

## Dumping/Backing up database from the MySQL Shell

```
mysqldump -u [username] -p [database-you-want-to-dump] > [file-path]
```

```
mysqldump -u root -p mydatabase > /home/myuser/database-dump.sql
```

## Copying Tables Between databases (on the same server)

```
INSERT INTO DB1.TABLE1 SELECT * FROM DB2.TABLE1;
```

Use **REPLACE INTO** instead  
of **INSERT INTO** to  
overwrite data.



# Merging Databases - Bonus

## Exporting a Table

```
mysqldump -u root -p --no-create-info db1 > db1.sql  
mysqldump -u root -p --no-create-info db2 > db2.sql  
mysqldump -u root -p --no-data db1 > schema.sql
```

## Create a new DB

```
mysql -uroot -p -Ddb3 < schema.sql  
mysql -uroot -p -Ddb3 < db1.sql  
mysql -uroot -p -Ddb3 < db2.sql
```



# Modification of the Database – Insertion

1. What would the DDL look like for the following where customer ID is to be auto-incremented?
2. How would you insert a customer record into the customer table of the following schema where customer ID is to be auto-incremented?

*branch(branch\_name, branch\_city, assets)*

*customer (ID, customer\_name, customer\_street, customer\_city)*

*loan (loan\_number, branch\_name, amount)*

*borrower (ID, loan\_number)*

*account (account\_number, branch\_name, balance )*

*depositor (ID, account\_number)*

## Customer DDL

```
create table customer (  
  ID int auto_increment,  
  customer-name varchar(20),  
  customer-street varchar(25),  
  customer-city varchar(25),  
  primary key(ID)  
)
```

# Modification of the Database – Programmatic Insertion Algorithm 1

Difficulty level: 

This algorithm requires the user to simply create insert statements as strings based on the format for each table. Then issue the insert statement using code.

Create a list of tables based on integrity constraints (i.e., foreign key order)

Example: tables = ["classroom", "department", "course", ...]

create insert statement templates based on each table structure:

Examples:

```
insertStatement1 = "insert into table1 values(%s, %s, ...);" //known as a parameterized query
```

```
insertStatement2 = "insert into table2 values(%s, %s, ...);"
```

For each csvFile:

```
csvFile = open(csvFile)
```

```
for each line in the csv file
```

```
create tuple from line as record //Python requires tuple whereas Java PreparedStatement
```

```
insert record into database table based on insertStatement template for this csvFile
```

Hey, Mr. DB here!  
See Python code example  
for inserting data on the  
last slide.



# Modification of the Database – Programmatic Insertion Algorithm 2

Difficulty level: 

This algorithm creates a DDL file based on a set of CSV Files. The DDL file will have statements of the form:

insert into ***table name*** values('value<sub>1</sub>', 'value<sub>2</sub>', ... 'value<sub>n</sub>');

Create a list of tables based on integrity constraints (i.e., foreign key order)

Example: tables = ["classroom", "department", "course", ...]

Determine a **file name** for the output DDL file

**fp** = open DDL file for writing

for each **table** in the list:

    let **csvFile** = open a csv file corresponding to the ***table name*** (eg. "classroom.csv")

        build a string from each record in the file as **record\_str** where each attribute except null is enclosed in single quotes

        set insertStatement as "insert into ***table name*** values(***record\_str***);"

        write insertStatement to **fp**

**Batch execute the DDL file to insert all data into database**

# Modification of the Database – Programmatic Insertion Algorithm 3

Difficulty level: 

This algorithm reads one or more CSV files, returns records corresponding to each CSV file as a list.  
The programmer will loop through each list and write the data to the database.

Create a list of tables based on integrity constraints (i.e., foreign key order)

Example: tables = ["classroom", "department", "course", ...]

Let **insertStatement** = "insert into %s values(%s);" //known as a parameterized query

for each **table** in the list:

let **csvFile** = open a csv file corresponding to the **table name** (eg. "classroom.csv")

for **record** in **csvFile**

build a string from each record in the file as **record\_str** where each attribute except null is enclosed in single quotes

let **wildcard** = "%s" or "?" based on programming language

let **wildcards** = **wildcard** \* (number of attributes in the record)

fill in the first %s in the **insertStatement** with the **table name**

fill in the second %s in the **insertStatement** with the value in the **wildcards** variable

issue the insert statement to the database // Note that substituting the wildcard values (i.e., ?,? ...) in Java

// will require some work based on the data type of each attribute.

# Python Code to Insert a Record into a Database

```
import mysql.connector
try:
    # Connecting to the Database
    conn = mysql.connector.connect(host='localhost', database='bankdb', user='root')
    # Variables may come from anywhere, a file, user input, etc
    customer_name = ""
    customer_street = ""
    customer_city = ""

    cursor = conn.cursor()
    sql_command = """INSERT INTO customer(customer_name, customer_street, customer_city) VALUES(%s, %s, %s);"""
    record = (customer_name, customer_street, customer_city) # a record is a tuple

    cursor.execute(sql_command, record) # the execute method will fill in each record value into the sql_command string

    # commit changes to the database
    conn.commit()

    # update successful message
    print("Customer data added successfully")

except mysql.connector.Error as error:
    conn.rollback()
    conn.close()
```

# Java Code to Insert Records into a Database

```
import java.sql.*;
import java.util.ArrayList;

public class SampleInsert {
    public void insertCustomerData(ArrayList<String[]> records){
        try {
            Connection conn = DriverManager.getConnection(DB, USER, PASS);
            String customerQuery = "INSERT INTO customer(customer_name, customer_street, customer_city) VALUES(?,?,?)";
            for (String[] record : records) {
                PreparedStatement preparedStatement = conn.prepareStatement(customerQuery);
                preparedStatement.setString(1, record[1]);
                preparedStatement.setString(2, record[2]);
                preparedStatement.setString(3, record[3]);
                preparedStatement.execute();
            }
            System.out.println("Inserted " + records.size() + " record(s) into customer table");
            conn.close();
        } catch (SQLException e) {
            System.err.println("An error occurred while trying to write data to the database");
            System.err.println(e.getMessage());
        }
    }
}
```

**End of Chapter**

